

Machine Learning Project

Plato Karageorgis

29/11/2021

A Quick Intro

As the report states “One thing that people regularly do is quantify how much of a particular activity they do, but they rarely quantify how well they do it” and that is what we are about to examine. We are going to create a model to predict the variable “classe” of our data. This variable pinpoints the way these people did their exercise, something that we can see thanks to the accelerometers that they wore.

```
train_in <- read.csv('./pml-training.csv', header=T)
```

To begin the analysis, we will assign our data to local variables.

```
train_data <- read.csv('C:/Users/Platon/Desktop/Coursera/pml-training.csv', header=T)
test_data <- read.csv('C:/Users/Platon/Desktop/Coursera/pml-testing.csv', header=T)
```

Since the test data will be used later in order to test, as the name states, our prediction model, we will start with the train data. Following the course’s instructions, we will use cross validation and separate the training data in partitions. The 75% will be used for training and the rest 25% will be responsible for the validation.

```
set.seed(33000)
inTrain <- createDataPartition(y= train_data$classe, p=0.7, list=FALSE)
training <- train_data[inTrain, ]
testing <- train_data[-inTrain, ]
```

Now, we have to start “playing” with our data. The main problem here is that we have a lot of NA values. This affects us mostly on our test set and not so much on our training set, although if we want to be precise, we should do it on both. There are multiples ways to go. We could erase all of them but then we might have a problem with the quantity of the remaining data. Also, we could erase some of them randomly and uniformly. I will choose the simpler solution and erase all of them since we have a lot of data.

```
nas <- sapply(names(test_data), function(x) all(is.na(test_data[,x])==TRUE))
nas_names <- names(nas)[nas==FALSE]
nas_names <- nas_names[-(1:7)]
nas_names <- nas_names[1:(length(nas_names)-1)]
fit <- trainControl(method='cv', number = 3)
```

Algorithms

I will apply 2 algorithms to check if my model is working correctly. I will use Random Forest and Boosting Trees. Firstly, for the random forest algorithm:

```
modelFitRF <- train(classe ~ ., data=training[, c('classe', nas_names)], trControl = fit, method='rf', n
```

And for the prediction:

```
predictFitRF <- predict(modelFitRF, testing, type = "raw")
```

And moving on the GBM algorithm:

```
modelFitGBM <- train(classe ~ ., data=training[, c('classe', nas_names)], trControl = fit, method='gbm', n
```

```
## Iter   TrainDeviance   ValidDeviance   StepSize   Improve
##      1         1.6094         nan         0.1000     0.1268
##      2         1.5240         nan         0.1000     0.0824
##      3         1.4673         nan         0.1000     0.0675
##      4         1.4225         nan         0.1000     0.0557
##      5         1.3872         nan         0.1000     0.0404
##      6         1.3594         nan         0.1000     0.0411
##      7         1.3325         nan         0.1000     0.0376
##      8         1.3072         nan         0.1000     0.0316
##      9         1.2852         nan         0.1000     0.0369
##     10         1.2615         nan         0.1000     0.0323
##     20         1.1040         nan         0.1000     0.0169
##     40         0.9313         nan         0.1000     0.0084
##     60         0.8211         nan         0.1000     0.0052
##     80         0.7388         nan         0.1000     0.0043
##    100         0.6748         nan         0.1000     0.0037
##    120         0.6215         nan         0.1000     0.0037
##    140         0.5780         nan         0.1000     0.0015
##    150         0.5589         nan         0.1000     0.0018
```

```
##
## Iter   TrainDeviance   ValidDeviance   StepSize   Improve
##      1         1.6094         nan         0.1000     0.1861
##      2         1.4903         nan         0.1000     0.1281
##      3         1.4064         nan         0.1000     0.1002
##      4         1.3411         nan         0.1000     0.0828
##      5         1.2875         nan         0.1000     0.0663
##      6         1.2437         nan         0.1000     0.0722
##      7         1.1972         nan         0.1000     0.0605
##      8         1.1585         nan         0.1000     0.0486
##      9         1.1256         nan         0.1000     0.0513
##     10         1.0930         nan         0.1000     0.0387
##     20         0.8917         nan         0.1000     0.0238
##     40         0.6746         nan         0.1000     0.0131
##     60         0.5426         nan         0.1000     0.0062
##     80         0.4561         nan         0.1000     0.0040
##    100         0.3892         nan         0.1000     0.0023
##    120         0.3404         nan         0.1000     0.0032
##    140         0.2997         nan         0.1000     0.0030
##    150         0.2815         nan         0.1000     0.0010
```

```
##
## Iter   TrainDeviance   ValidDeviance   StepSize   Improve
##      1         1.6094         nan         0.1000     0.2245
```

##	2	1.4646	nan	0.1000	0.1610
##	3	1.3614	nan	0.1000	0.1184
##	4	1.2865	nan	0.1000	0.1044
##	5	1.2183	nan	0.1000	0.0841
##	6	1.1629	nan	0.1000	0.0705
##	7	1.1174	nan	0.1000	0.0700
##	8	1.0733	nan	0.1000	0.0595
##	9	1.0354	nan	0.1000	0.0622
##	10	0.9973	nan	0.1000	0.0592
##	20	0.7583	nan	0.1000	0.0244
##	40	0.5294	nan	0.1000	0.0119
##	60	0.4001	nan	0.1000	0.0077
##	80	0.3167	nan	0.1000	0.0054
##	100	0.2580	nan	0.1000	0.0032
##	120	0.2140	nan	0.1000	0.0023
##	140	0.1827	nan	0.1000	0.0014
##	150	0.1691	nan	0.1000	0.0010
##					
##	Iter	TrainDeviance	ValidDeviance	StepSize	Improve
##	1	1.6094	nan	0.1000	0.1275
##	2	1.5226	nan	0.1000	0.0868
##	3	1.4644	nan	0.1000	0.0648
##	4	1.4205	nan	0.1000	0.0525
##	5	1.3852	nan	0.1000	0.0406
##	6	1.3572	nan	0.1000	0.0446
##	7	1.3284	nan	0.1000	0.0394
##	8	1.3035	nan	0.1000	0.0334
##	9	1.2816	nan	0.1000	0.0335
##	10	1.2591	nan	0.1000	0.0323
##	20	1.1001	nan	0.1000	0.0171
##	40	0.9313	nan	0.1000	0.0083
##	60	0.8206	nan	0.1000	0.0050
##	80	0.7405	nan	0.1000	0.0033
##	100	0.6787	nan	0.1000	0.0036
##	120	0.6264	nan	0.1000	0.0036
##	140	0.5833	nan	0.1000	0.0015
##	150	0.5633	nan	0.1000	0.0026
##					
##	Iter	TrainDeviance	ValidDeviance	StepSize	Improve
##	1	1.6094	nan	0.1000	0.1891
##	2	1.4883	nan	0.1000	0.1263
##	3	1.4061	nan	0.1000	0.1020
##	4	1.3401	nan	0.1000	0.0872
##	5	1.2840	nan	0.1000	0.0790
##	6	1.2347	nan	0.1000	0.0629
##	7	1.1948	nan	0.1000	0.0577
##	8	1.1576	nan	0.1000	0.0491
##	9	1.1262	nan	0.1000	0.0416
##	10	1.0984	nan	0.1000	0.0434
##	20	0.8974	nan	0.1000	0.0201
##	40	0.6839	nan	0.1000	0.0097
##	60	0.5569	nan	0.1000	0.0054
##	80	0.4666	nan	0.1000	0.0056
##	100	0.3984	nan	0.1000	0.0045

##	120	0.3455	nan	0.1000	0.0033
##	140	0.3028	nan	0.1000	0.0021
##	150	0.2831	nan	0.1000	0.0015
##					
##	Iter	TrainDeviance	ValidDeviance	StepSize	Improve
##	1	1.6094	nan	0.1000	0.2358
##	2	1.4612	nan	0.1000	0.1638
##	3	1.3563	nan	0.1000	0.1291
##	4	1.2751	nan	0.1000	0.1042
##	5	1.2091	nan	0.1000	0.0792
##	6	1.1576	nan	0.1000	0.0698
##	7	1.1126	nan	0.1000	0.0792
##	8	1.0628	nan	0.1000	0.0577
##	9	1.0250	nan	0.1000	0.0539
##	10	0.9896	nan	0.1000	0.0512
##	20	0.7584	nan	0.1000	0.0248
##	40	0.5329	nan	0.1000	0.0190
##	60	0.4004	nan	0.1000	0.0112
##	80	0.3118	nan	0.1000	0.0034
##	100	0.2579	nan	0.1000	0.0022
##	120	0.2154	nan	0.1000	0.0011
##	140	0.1842	nan	0.1000	0.0010
##	150	0.1709	nan	0.1000	0.0022
##					
##	Iter	TrainDeviance	ValidDeviance	StepSize	Improve
##	1	1.6094	nan	0.1000	0.1285
##	2	1.5234	nan	0.1000	0.0906
##	3	1.4655	nan	0.1000	0.0654
##	4	1.4211	nan	0.1000	0.0523
##	5	1.3852	nan	0.1000	0.0436
##	6	1.3556	nan	0.1000	0.0434
##	7	1.3269	nan	0.1000	0.0379
##	8	1.3029	nan	0.1000	0.0350
##	9	1.2806	nan	0.1000	0.0352
##	10	1.2589	nan	0.1000	0.0323
##	20	1.1031	nan	0.1000	0.0189
##	40	0.9297	nan	0.1000	0.0099
##	60	0.8189	nan	0.1000	0.0061
##	80	0.7405	nan	0.1000	0.0029
##	100	0.6758	nan	0.1000	0.0026
##	120	0.6234	nan	0.1000	0.0029
##	140	0.5793	nan	0.1000	0.0019
##	150	0.5605	nan	0.1000	0.0023
##					
##	Iter	TrainDeviance	ValidDeviance	StepSize	Improve
##	1	1.6094	nan	0.1000	0.1847
##	2	1.4892	nan	0.1000	0.1293
##	3	1.4059	nan	0.1000	0.1045
##	4	1.3388	nan	0.1000	0.0861
##	5	1.2823	nan	0.1000	0.0667
##	6	1.2379	nan	0.1000	0.0590
##	7	1.1993	nan	0.1000	0.0639
##	8	1.1596	nan	0.1000	0.0575
##	9	1.1230	nan	0.1000	0.0445

##	10	1.0937	nan	0.1000	0.0455
##	20	0.8870	nan	0.1000	0.0227
##	40	0.6712	nan	0.1000	0.0077
##	60	0.5475	nan	0.1000	0.0064
##	80	0.4628	nan	0.1000	0.0058
##	100	0.3951	nan	0.1000	0.0028
##	120	0.3412	nan	0.1000	0.0016
##	140	0.3024	nan	0.1000	0.0025
##	150	0.2838	nan	0.1000	0.0014
##					
##	Iter	TrainDeviance	ValidDeviance	StepSize	Improve
##	1	1.6094	nan	0.1000	0.2385
##	2	1.4580	nan	0.1000	0.1563
##	3	1.3587	nan	0.1000	0.1208
##	4	1.2822	nan	0.1000	0.0989
##	5	1.2175	nan	0.1000	0.0878
##	6	1.1617	nan	0.1000	0.0737
##	7	1.1139	nan	0.1000	0.0663
##	8	1.0706	nan	0.1000	0.0644
##	9	1.0289	nan	0.1000	0.0468
##	10	0.9989	nan	0.1000	0.0539
##	20	0.7576	nan	0.1000	0.0222
##	40	0.5220	nan	0.1000	0.0136
##	60	0.3952	nan	0.1000	0.0056
##	80	0.3151	nan	0.1000	0.0045
##	100	0.2586	nan	0.1000	0.0026
##	120	0.2148	nan	0.1000	0.0022
##	140	0.1820	nan	0.1000	0.0013
##	150	0.1685	nan	0.1000	0.0013
##					
##	Iter	TrainDeviance	ValidDeviance	StepSize	Improve
##	1	1.6094	nan	0.1000	0.2301
##	2	1.4629	nan	0.1000	0.1600
##	3	1.3615	nan	0.1000	0.1155
##	4	1.2879	nan	0.1000	0.1024
##	5	1.2238	nan	0.1000	0.0885
##	6	1.1680	nan	0.1000	0.0733
##	7	1.1218	nan	0.1000	0.0715
##	8	1.0776	nan	0.1000	0.0676
##	9	1.0346	nan	0.1000	0.0632
##	10	0.9949	nan	0.1000	0.0508
##	20	0.7650	nan	0.1000	0.0295
##	40	0.5370	nan	0.1000	0.0131
##	60	0.4086	nan	0.1000	0.0100
##	80	0.3281	nan	0.1000	0.0044
##	100	0.2683	nan	0.1000	0.0023
##	120	0.2237	nan	0.1000	0.0026
##	140	0.1894	nan	0.1000	0.0016
##	150	0.1757	nan	0.1000	0.0016

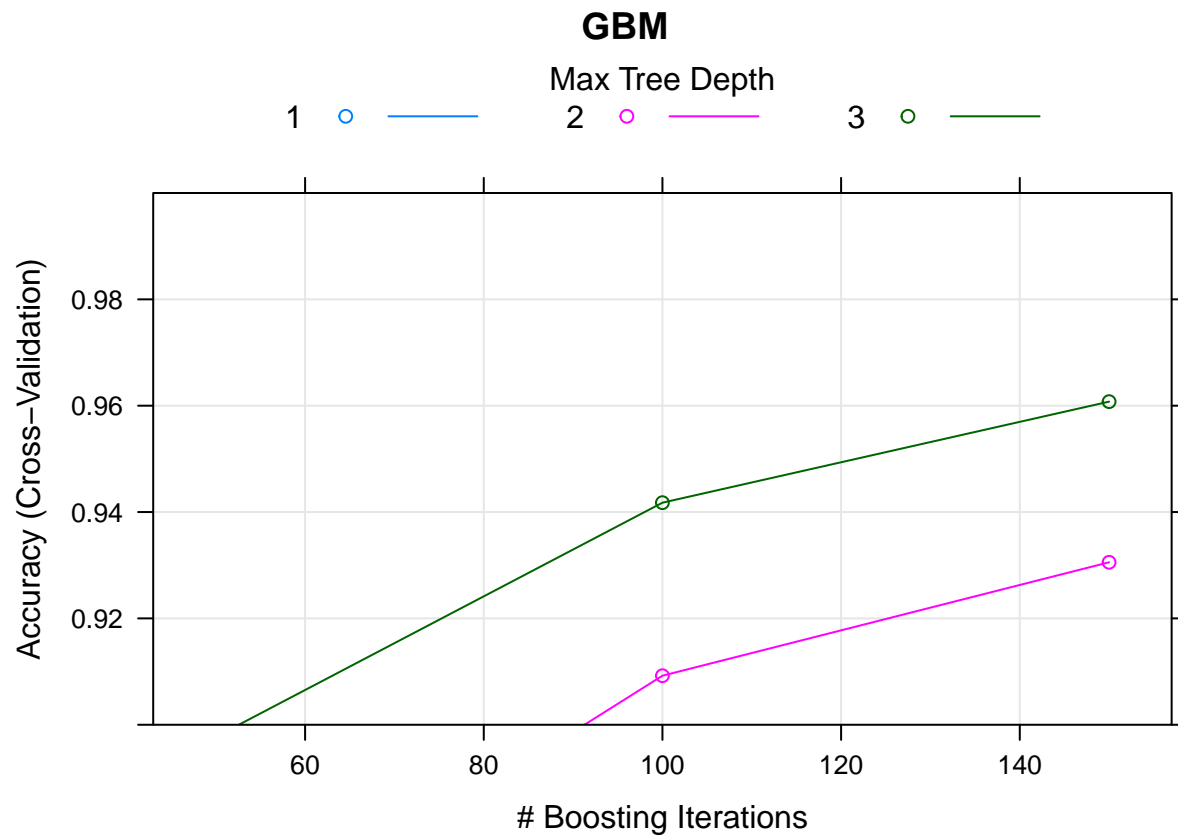
And for the prediction:

```
predictFitGBM <- predict(modelFitGBM, testing, type = "raw")
```


Plots

Moving on, I will do some very simple plots with the plot() function in order to visualize the accuracies.

```
plot(modelFitGBM, ylim = c(0.9, 1), main= "GBM")
```



```
plot(modelFitRF, ylim = c(0.9, 1),main= "Random Forest")
```



Results

Moreover, I am creating a confusion matrix exactly like the ones shown in the lectures in order to see the exact values of the accuracies.

```
confusionGBM <- confusionMatrix(predictFitGBM, as.factor(testing$classe))
confusionRF <- confusionMatrix(predictFitRF, as.factor(testing$classe))

accuracy <- data.frame(Model = c('RF', 'GBM'), Accuracy = rbind(confusionRF$overall[1], confusionGBM$overall[1]))
print(accuracy)
```

```
##   Model Accuracy
## 1    RF 0.9935429
## 2    GBM 0.9639762
```

Final Predictions

Last but not least, the 20 predictions that the exercise requests. I will use Random Forest since it was better.

```
predictions <- predict(modelFitRF, newdata=test_data)
final <- data.frame(problem_id= test_data$problem_id, predicted=predictions)
print(final)
```

##	problem_id	predicted
## 1	1	B
## 2	2	A
## 3	3	B
## 4	4	A
## 5	5	A
## 6	6	E
## 7	7	D
## 8	8	B
## 9	9	A
## 10	10	A
## 11	11	B
## 12	12	C
## 13	13	B
## 14	14	A
## 15	15	E
## 16	16	E
## 17	17	A
## 18	18	B
## 19	19	B
## 20	20	B

Conclusion

We can clearly see that the random forest algorithm is more than capable of predicting the right values from the plots but even more clearly on our matrix. The GBM model takes is also running quite impressively. You can see the results on the .html file.